

In [1]:

```
1 !wget --header="Host: storage.googleapis.com" --header="User-Agent: Mozilla/5.0 (Windows
```

```
--2021-04-30 15:08:57-- https://storage.googleapis.com/kagglesdsdata/competitions/7634/46676/train.7z?GoogleAccessId=web-data@kaggle-161607.iam.gserviceaccount.com&Expires=1620017510&Signature=XIU9aJLJcDgoFuUsl31Iw762mbcQzn1A8sdD3aKfT7le7lP%2BRNALcMQwsKf2aZPu49EIUxkHjqyT8m2Cdkv74rixFFBXXUefuA70e1EcTl7BCUEC5mj3wbw8Cc22c6SaJRLqjeQMLn%2B3H0Gy74h7IvtZLKDsRvIJxHS7jlA8ukFLr3%2FjuoftTmRRCyF7b6AO3mUWMNSH%2Fn3I83CyvKSg1RJg801JLvCS8mZ96DeMOqPQRbcvWVyHHJZxFsLI90yhOZIqd0090izi6%2Bn0%2B%2BiL%2BIHa7QjjkvxhlWiWU1MeBw2De%2F1VjTdnZnw%2BCwQS6jrglspcunmjZN65gwpGwgvuQ%3D%3D&response-content-disposition=attachment%3B+filename%3Dtrain.7z (https://storage.googleapis.com/kagglesdsdata/competitions/7634/46676/train.7z?GoogleAccessId=web-data@kaggle-161607.iam.gserviceaccount.com&Expires=1620017510&Signature=XIU9aJLJcDgoFuUsl31Iw762mbcQzn1A8sdD3aKfT7le7lP%2BRNALcMQwsKf2aZPu49EIUxkHjqyT8m2Cdkv74rixFFBXXUefuA70e1EcTl7BCUEC5mj3wbw8Cc22c6SaJRLqjeQMLn%2B3H0Gy74h7IvtZLKDsRvIJxHS7jlA8ukFLr3%2FjuoftTmRRCyF7b6AO3mUWMNSH%2Fn3I83CyvKSg1RJg801JLvCS8mZ96DeMOqPQRbcvWVyHHJZxFsLI90yhOZIqd0090izi6%2Bn0%2B%2BiL%2BIHa7QjjkvxhlWiWU1MeBw2De%2F1VjTdnZnw%2BCwQS6jrglspcunmjZN65gwpGwgvuQ%3D%3D&response-content-disposition=attachment%3B+filename%3Dtrain.7z)
Resolving storage.googleapis.com (storage.googleapis.com)... 74.125.20.128, 74.125.199.128, 74.125.142.128, ...
Connecting to storage.googleapis.com (storage.googleapis.com)|74.125.20.128|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 1121103842 (1.0G) [application/x-7z-compressed]
Saving to: 'train.7z'

train.7z          100%[=====>]    1.04G   64.2MB/s   in 19s

2021-04-30 15:09:16 (56.7 MB/s) - 'train.7z' saved [1121103842/1121103842]
```

In [2]:

```
1 !apt-get install p7zip-full
2 !p7zip -d train.7z
```

```
Reading package lists... Done
Building dependency tree
Reading state information... Done
p7zip-full is already the newest version (16.02+dfsg-6).
The following package was automatically installed and is no longer require
d:
  libnvidia-common-460
Use 'apt autoremove' to remove it.
0 upgraded, 0 newly installed, 0 to remove and 34 not upgraded.

7-Zip (a) [64] 16.02 : Copyright (c) 1999-2016 Igor Pavlov : 2016-05-21
p7zip Version 16.02 (locale=en_US.UTF-8,Utf16=on,HugeFiles=on,64 bits,2 CP
Us Intel(R) Xeon(R) CPU @ 2.20GHz (406F0),ASM,AES-NI)

Scanning the drive for archives:
  0M Sca          1 file, 1121103842 bytes (1070 MiB)

Extracting archive: train.7z
--
~ .. . . .
```

In [3]:

```
1 from google.colab import drive
2 drive.mount('/content/drive')
```

Mounted at /content/drive

In [4]:

```

1 import os
2 import pandas as pd
3 import numpy as np
4 import csv
5 #!pip install tqdm
6 from tqdm import tqdm_notebook as tqdm
7 from collections import Counter
8 import matplotlib.pyplot as plt
9 import seaborn as sns
10 import warnings
11 warnings.filterwarnings("ignore")
12 from collections import Counter
13 import librosa
14 import IPython.display as ipd
15 import librosa.display
16 from scipy.io import wavfile
17 from scipy.fftpack import fft
18 from scipy import signal
19
20 from sklearn.model_selection import train_test_split
21 from tensorflow.keras.callbacks import ReduceLROnPlateau, TensorBoard
22 from sklearn.utils import shuffle
23 from tensorflow.keras.layers import Input, LSTM, Dense, AveragePooling2D, TimeDistrib
24 from tensorflow.keras.models import Model
25 import tensorflow as tf
26
27 from sklearn.preprocessing import LabelEncoder
28
29 from sklearn.preprocessing import OneHotEncoder
30 import pickle
31
32 from sklearn.metrics import classification_report
33
34 os.chdir('/content/train/audio')

```

In [5]:

```

1 data = pd.read_csv('/content/drive/MyDrive/cs-2/data/file_name+label.csv')
2 data = data[data['label'] != '_background_noise_']
3 print(data.shape)
4 #data.head(3)

```

(64721, 2)

In [7]:

```
1 data.head(3)
```

Out[7]:

	file_name	label
0	five/c518d1b1_nohash_1.wav	five
1	five/f17be97f_nohash_1.wav	five
2	five/20d779bf_nohash_0.wav	five

In []:

```
1 data = shuffle(data)
```

In []:

```
1 # get all raw data
2 raw_data = []
3 for i in tqdm(data['file_name']):
4     signal, rate = librosa.load(i, sr=16000)
5     #signal = List(signal)
6     raw_data.append(signal)
```

```
HBox(children=(FloatProgress(value=0.0, max=64721.0), HTML(value='')))
```

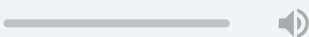
Plot Raw data

In []:

```
1 idx = 100
2 samples = raw_data[idx]
3 sample_rate = 16000
4 print('Label : ', data['label'][idx:idx+1].values[0])
5 ipd.Audio(samples, rate = sample_rate)
```


Label : one

Out[9]:

▶ 0:00 / 0:00 

In []:

```
1 sns.set_theme()
2 fig = plt.figure(figsize=(14, 8))
3 ax = fig.add_subplot(211)
4
5 ax.set_title('Raw wave of '+'One')
6 ax.set_xlabel('Time(second)')
7 ax.set_ylabel('Amplitude')
8 ax.plot(np.linspace(0, sample_rate/len(samples), sample_rate), samples)
9 plt.show()
```



In []:

```

1 def plot_signals(signals):
2     fig , ax = plt.subplots(nrows = 6 , ncols = 5 , figsize = (20,15))
3     i = 0
4     for x in range(6):
5         for y in range(5):
6             ax[x,y].set_title(list(signals.keys())[i])
7             ax[x,y].plot(list(signals.values())[i])
8             ax[x,y].get_xaxis().set_visible(False)
9             ax[x,y].get_yaxis().set_visible(False)
10            i += 1

```

In []:

```

1 # store all signal in dic
2 signals = {}
3 labels = np.unique(data['label'])
4
5 # get all signal array except backgroud noise
6 for name in labels:
7     file = data[data['label'] == name][:1]
8
9     signal , rate = librosa.load(file['file_name'].values[0])
10
11     signals[name] = signal

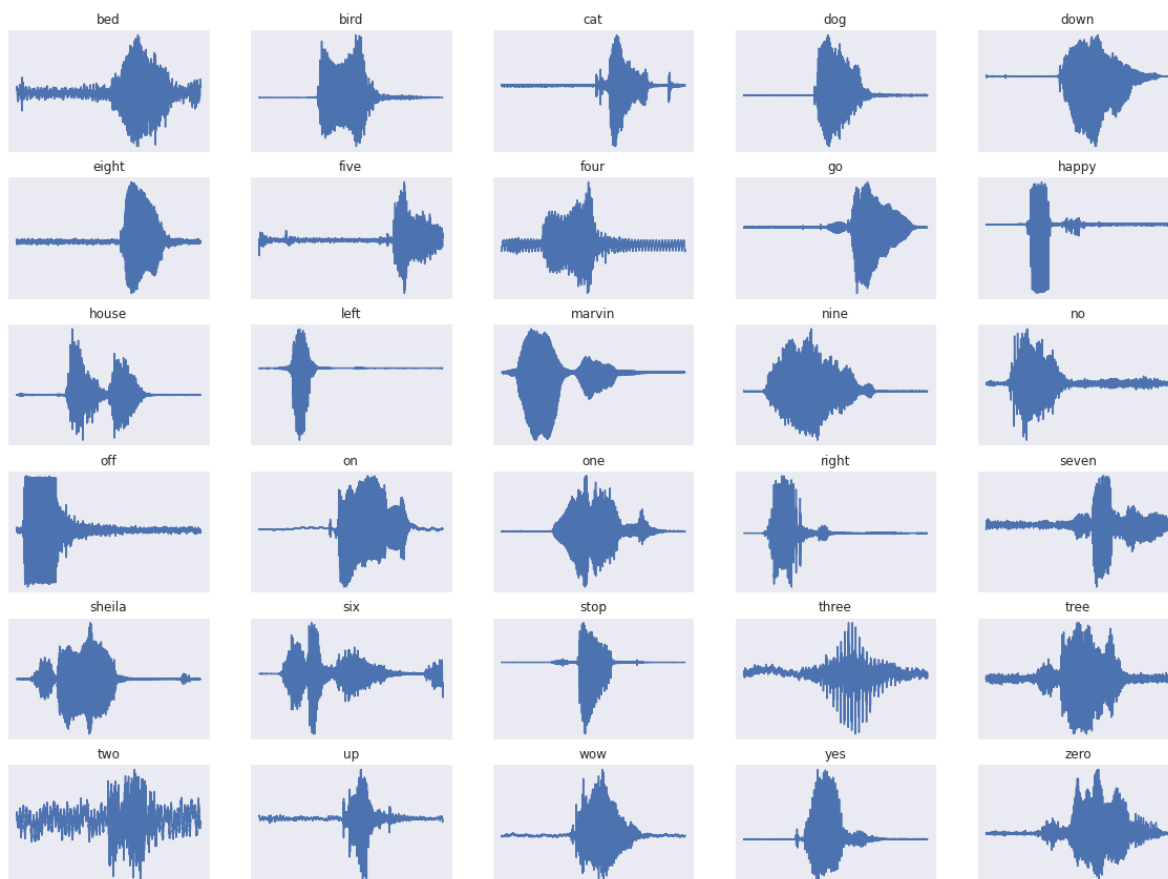
```

In []:

```

1 plot_signals(signals)
2 plt.show()

```



As we see in plot many sound wave is nearly zero at starting. So we trim that silence at starting.

Padding

In []:

```

1  ''' padding - make all file same length(16000) '''
2  # if len of audio is < 16000 we add zeros.
3  # if len of audio is > 16000 we truncate.
4
5  def padding(data, size):
6      if len(data) >= size:
7          arr = data[:size]
8      else:
9          arr = np.zeros(size)
10         arr[:len(data)] = data
11     return arr

```

In []:

```

1  pad_seq = []
2  for i in tqdm(raw_data):
3      i = padding(i,16000)
4      pad_seq.append(i)

```

HBox(children=(FloatProgress(value=0.0, max=64721.0), HTML(value='')))

In []:

```

1  #data['raw_data'] = raw_data
2  data['pad_seq'] = pad_seq
3  del raw_data, pad_seq

```

train test split

In []:

```

1  data = shuffle(data,random_state=33)
2
3  y = data['label']
4  data = data.drop('label',axis = 1)
5
6  x_train , x_vali, y_train, y_vali = train_test_split(data['pad_seq'].values, y,test_si

```

In []:

```

1  print(x_train.shape, y_train.shape)
2  print(x_vali.shape, y_vali.shape)

```

```

(51776,) (51776,)
(12945,) (12945,)

```

In []:

```

1 #for error analysis
2 np.save('/content/train_pad.npy', x_train)
3 y_train.to_csv('/content/y_train.csv', index = False)
4
5 np.save('/content/vali_pad.npy', x_vali)
6 y_vali.to_csv('/content/y_vali.csv', index = False)

```

In [7]:

```
1 del raw_train , raw_vali, raw_test
```

load data

In []:

```

1 x_train = np.load('/content/pad_train.npy', allow_pickle=True)
2 x_vali = np.load('/content/pad_vali.npy', allow_pickle=True)
3
4 y_train = pd.read_csv('/content/y_train.csv')
5 y_vali = pd.read_csv('/content/y_vali.csv')

```

In []:

```

1 print(x_train.shape, y_train.shape)
2 print(x_vali.shape, y_vali.shape )

```

```

(51776,) (51776,)
(12945,) (12945,)

```

log_spectrogram

In []:

```

1 def log_spectrogram(audio, sample_rate, window_size=20,
2                     step_size=10, eps=1e-10):
3     nperseg = int(round(window_size * sample_rate / 1e3))
4     noverlap = int(round(step_size * sample_rate / 1e3))
5     freqs, times, spec = signal.spectrogram(audio,
6                                             fs=sample_rate,
7                                             window='hann',
8                                             nperseg = nperseg,
9                                             noverlap = noverlap,
10                                            detrend=False)
11     return freqs, times, np.log(spec.T.astype(np.float32) + eps)

```

In []:

```
1 from scipy import signal
```

In []:

```
1 train_data = []
2 for i in tqdm(x_train):
3     _,_,spec = log_spectrogram(i,16000)
4     train_data.append(spec)
5 train_data = np.array(train_data)
6 train_data.shape
```

HBox(children=(FloatProgress(value=0.0, max=51776.0), HTML(value='')))

Out[25]:

(51776, 99, 161)

In []:

```
1 np.save('/content/log_train_data.npy', train_data)
2 del train_data
```

In []:

```
1 vali_data = []
2 for i in tqdm(x_vali):
3     _,_,spec = log_spectrogram(i,16000)
4     vali_data.append(spec)
5 vali_data = np.array(vali_data)
6 vali_data.shape
```

HBox(children=(FloatProgress(value=0.0, max=12945.0), HTML(value='')))

Out[28]:

(12945, 99, 161)

In []:

```
1 np.save('/content/log_vali_data.npy', vali_data)
```

Load Final Files

In []:

```
1 train_data = np.load('/content/log_train_data.npy', allow_pickle=True)
2 vali_data = np.load('/content/log_vali_data.npy', allow_pickle=True)
3
4 y_train = pd.read_csv('/content/y_train.csv')
5 y_vali = pd.read_csv('/content/y_vali.csv')
```


In []:

```
1 # reshape because we use conv2d Here
2 train_data = train_data.reshape(tuple(list(train_data.shape) + [1] ))
3 train_data.shape
```

Out[3]:

(51776, 99, 161, 1)

In []:

```
1 vali_data = vali_data.reshape(tuple(list(vali_data.shape) + [1] ))
2 vali_data.shape
```

Out[4]:

(12945, 99, 161, 1)

In []:

```
1 total = y_train.append(y_vali)
2 print(len(total.label.unique()))
3 total.label.unique()
```

30

Out[5]:

```
array(['go', 'left', 'tree', 'eight', 'yes', 'cat', 'two', 'up', 'down',
      'no', 'dog', 'sheila', 'happy', 'off', 'on', 'nine', 'three',
      'zero', 'bird', 'seven', 'house', 'bed', 'six', 'stop', 'one',
      'marvin', 'right', 'wow', 'five', 'four'], dtype=object)
```

encoder

In []:

```
1 encoder = LabelEncoder()
2 encoder.fit(total)
3 y_train = encoder.transform(y_train)
4 y_vali = encoder.transform(y_vali)
5 #y_test = encoder.transform(y_test)
```

In []:

```
1 y_train.shape , y_vali.shape
```

Out[7]:

((51776,), (12945,))

Model

In []:

```
1 model_7 = tf.keras.Sequential()
2
3 model_7.add(Input(shape=(99,161,1)))
4 model_7.add(BatchNormalization())
5 model_7.add(Conv2D(8,kernel_size=2,activation='relu',kernel_regularizer=tf.keras.regul
6 model_7.add(Conv2D(8,kernel_size=3,activation='relu',kernel_regularizer=tf.keras.regul
7 model_7.add(MaxPooling2D(pool_size=(2,2)))
8 model_7.add(Dropout(rate=0.2))
9
10 model_7.add(Conv2D(16,kernel_size=5,activation='relu',kernel_regularizer=tf.keras.regu
11 model_7.add(Conv2D(16,kernel_size=6,activation='relu',kernel_regularizer=tf.keras.regu
12 model_7.add(MaxPooling2D(pool_size=(2,2)))
13 model_7.add(Dropout(rate=0.2))
14
15 model_7.add(Conv2D(16,kernel_size=6,activation='relu',kernel_regularizer=tf.keras.regu
16 model_7.add(MaxPooling2D(pool_size=(2,2)))
17
18 model_7.add(Flatten())
19 model_7.add(Dense(256, activation='relu'))
20 model_7.add(BatchNormalization())
21 model_7.add(Dense(128, activation='relu'))
22 model_7.add(BatchNormalization())
23 model_7.add(Dense(30, activation='softmax'))
```

In []:

```
1 model_7.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
batch_normalization (Batch Normalization)	(None, 99, 161, 1)	4
conv2d (Conv2D)	(None, 98, 160, 8)	40
conv2d_1 (Conv2D)	(None, 96, 158, 8)	584
max_pooling2d (MaxPooling2D)	(None, 48, 79, 8)	0
dropout (Dropout)	(None, 48, 79, 8)	0
conv2d_2 (Conv2D)	(None, 44, 75, 16)	3216
conv2d_3 (Conv2D)	(None, 39, 70, 16)	9232
max_pooling2d_1 (MaxPooling2D)	(None, 19, 35, 16)	0
dropout_1 (Dropout)	(None, 19, 35, 16)	0
conv2d_4 (Conv2D)	(None, 14, 30, 16)	9232
max_pooling2d_2 (MaxPooling2D)	(None, 7, 15, 16)	0
flatten (Flatten)	(None, 1680)	0
dense (Dense)	(None, 256)	430336
batch_normalization_1 (Batch Normalization)	(None, 256)	1024
dense_1 (Dense)	(None, 128)	32896
batch_normalization_2 (Batch Normalization)	(None, 128)	512
dense_2 (Dense)	(None, 30)	3870
=====		
Total params: 490,946		
Trainable params: 490,176		
Non-trainable params: 770		
=====		

In []:

```
1 tensorboard_callback = TensorBoard(log_dir='/content/model-8-3',write_grads=True,histo
2
3 def scheduler(epoch, lr):
4     if epoch < 5:
5         return lr
6     else:
7         return lr * 0.9
8
9 lrs = tf.keras.callbacks.LearningRateScheduler(scheduler,verbose=1)
10 callbacks = [tensorboard_callback,lrs]
```

WARNING:tensorflow:`write_grads` will be ignored in TensorFlow 2.0 for the `TensorBoard` Callback.

In []:

```
1 optimizer = tf.keras.optimizers.Adam(learning_rate=0.001)
2 opt = tf.keras.optimizers.SGD(learning_rate=0.001, momentum=0.9)
3 model_7.compile(optimizer=optimizer,loss='sparse_categorical_crossentropy', metrics =
```

In []:

```

1 history = model_7.fit( train_data ,y_train,
2                         epochs = 20,
3                         validation_data = ( vali_data ,y_vali) ,
4                         callbacks = callbacks)

```

Epoch 1/20

Epoch 00001: LearningRateScheduler reducing learning rate to 0.00100000004
74974513.

1618/1618 [=====] - 57s 15ms/step - loss: 2.2587
- accuracy: 0.3725 - val_loss: 0.6978 - val_accuracy: 0.8088

Epoch 2/20

Epoch 00002: LearningRateScheduler reducing learning rate to 0.00100000004
74974513.

1618/1618 [=====] - 22s 14ms/step - loss: 0.7165
- accuracy: 0.8044 - val_loss: 0.4898 - val_accuracy: 0.8742

Epoch 3/20

Epoch 00003: LearningRateScheduler reducing learning rate to 0.00100000004
74974513.

1618/1618 [=====] - 22s 14ms/step - loss: 0.5403
- accuracy: 0.8530 - val_loss: 0.4718 - val_accuracy: 0.8774

Epoch 4/20

Epoch 00004: LearningRateScheduler reducing learning rate to 0.00100000004
74974513.

1618/1618 [=====] - 22s 14ms/step - loss: 0.4637
- accuracy: 0.8787 - val_loss: 0.4215 - val_accuracy: 0.9003

Epoch 5/20

Epoch 00005: LearningRateScheduler reducing learning rate to 0.00100000004
74974513.

1618/1618 [=====] - 23s 14ms/step - loss: 0.4171
- accuracy: 0.8936 - val_loss: 0.4843 - val_accuracy: 0.8793

Epoch 6/20

Epoch 00006: LearningRateScheduler reducing learning rate to 0.00090000004
27477062.

1618/1618 [=====] - 23s 14ms/step - loss: 0.3848
- accuracy: 0.9029 - val_loss: 0.4282 - val_accuracy: 0.8949

Epoch 7/20

Epoch 00007: LearningRateScheduler reducing learning rate to 0.00081000003
84729356.

1618/1618 [=====] - 23s 14ms/step - loss: 0.3505
- accuracy: 0.9135 - val_loss: 0.3838 - val_accuracy: 0.9080

Epoch 8/20

Epoch 00008: LearningRateScheduler reducing learning rate to 0.00072900005
03417104.

1618/1618 [=====] - 23s 14ms/step - loss: 0.3164
- accuracy: 0.9237 - val_loss: 0.3450 - val_accuracy: 0.9190

Epoch 9/20

Epoch 00009: LearningRateScheduler reducing learning rate to 0.00065610007
15009868.

1618/1618 [=====] - 23s 14ms/step - loss: 0.2878
- accuracy: 0.9297 - val_loss: 0.3622 - val_accuracy: 0.9130

Epoch 10/20

Epoch 00010: LearningRateScheduler reducing learning rate to 0.0005904900433961303.

1618/1618 [=====] - 23s 14ms/step - loss: 0.2725
- accuracy: 0.9353 - val_loss: 0.3316 - val_accuracy: 0.9224

Epoch 11/20

Epoch 00011: LearningRateScheduler reducing learning rate to 0.0005314410547725857.

1618/1618 [=====] - 23s 14ms/step - loss: 0.2551
- accuracy: 0.9374 - val_loss: 0.3238 - val_accuracy: 0.9214

Epoch 12/20

Epoch 00012: LearningRateScheduler reducing learning rate to 0.00047829695977270604.

1618/1618 [=====] - 23s 14ms/step - loss: 0.2375
- accuracy: 0.9422 - val_loss: 0.3242 - val_accuracy: 0.9236

Epoch 13/20

Epoch 00013: LearningRateScheduler reducing learning rate to 0.0004304672533180565.

1618/1618 [=====] - 23s 14ms/step - loss: 0.2252
- accuracy: 0.9452 - val_loss: 0.3025 - val_accuracy: 0.9306

Epoch 14/20

Epoch 00014: LearningRateScheduler reducing learning rate to 0.00038742052274756136.

1618/1618 [=====] - 23s 14ms/step - loss: 0.2040
- accuracy: 0.9510 - val_loss: 0.3108 - val_accuracy: 0.9251

Epoch 15/20

Epoch 00015: LearningRateScheduler reducing learning rate to 0.0003486784757114947.

1618/1618 [=====] - 23s 14ms/step - loss: 0.1973
- accuracy: 0.9539 - val_loss: 0.3004 - val_accuracy: 0.9285

Epoch 16/20

Epoch 00016: LearningRateScheduler reducing learning rate to 0.00031381062290165574.

1618/1618 [=====] - 23s 14ms/step - loss: 0.1857
- accuracy: 0.9560 - val_loss: 0.2937 - val_accuracy: 0.9301

Epoch 17/20

Epoch 00017: LearningRateScheduler reducing learning rate to 0.0002824295632308349.

1618/1618 [=====] - 23s 14ms/step - loss: 0.1723
- accuracy: 0.9610 - val_loss: 0.2984 - val_accuracy: 0.9244

Epoch 18/20

Epoch 00018: LearningRateScheduler reducing learning rate to 0.00025418660952709616.

1618/1618 [=====] - 23s 14ms/step - loss: 0.1718
- accuracy: 0.9601 - val_loss: 0.2963 - val_accuracy: 0.9293

Epoch 19/20

Epoch 00019: LearningRateScheduler reducing learning rate to 0.00022876793809700757.

1618/1618 [=====] - 23s 14ms/step - loss: 0.1584
- accuracy: 0.9645 - val_loss: 0.2992 - val_accuracy: 0.9294

Epoch 20/20

Epoch 00020: LearningRateScheduler reducing learning rate to 0.00020589114428730683.
 1618/1618 [=====] - 23s 14ms/step - loss: 0.1578
 - accuracy: 0.9646 - val_loss: 0.2910 - val_accuracy: 0.9303

In []:

```
1 model_7.save('/content/drive/MyDrive/cs-2/final_model.h5')
2 np.save('/content/drive/MyDrive/cs-2/final_model_classes.npy', encoder.classes_)
```

In []:

```
1 #model 9
2 %reload_ext tensorboard
3 %tensorboard --logdir /content/model-8-3
4
```

<IPython.core.display.Javascript object>

Prediction on Validation data

In []:

```
1 predictions = model_7.predict_classes(vali_data, verbose=False)
2
3 prediction_labels = encoder.inverse_transform(predictions)
4
5 y = encoder.inverse_transform(y_vali)
```

In []:

```
1 df = pd.DataFrame()
2 df['prediction'] = prediction_labels
3 df['true_labels'] = y
```

In []:

```
1 df.head()
```

Out[20]:

	prediction	true_labels
0	four	four
1	on	on
2	nine	nine
3	marvin	marvin
4	stop	stop

In []:

```
1 df.to_csv('/content/df.csv', index=False)
```

In []:

```
1 # load validation pad data for plotting
2 pad_vali = np.load('/content/vali_pad.npy', allow_pickle=True)
```

In []:

```
1 df['pad_data'] = pad_vali
```

Create Dataframe with misclassified points

In []:

```
1 FP = df[df['prediction'] != df['true_labels']]
2 FP.shape
```

Out[29]:

(902, 3)

In []:

```
1 FP.head()
```

Out[30]:

	prediction	true_labels	pad_data
29	on	one	[-3.0517578e-05, -9.1552734e-05, -9.1552734e-0...
53	dog	down	[0.0049743652, 0.0056762695, 0.005432129, 0.00...
61	nine	marvin	[-0.00021362305, -0.0014343262, -0.002105713, ...
62	up	cat	[0.0, 0.0, 0.0, 3.0517578e-05, 0.0, 0.0, 0.0, ...
83	nine	marvin	[-3.0517578e-05, -3.0517578e-05, 6.1035156e-05...

Plot Misclassified Points

In []:

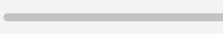

```


1  sns.set_theme()
2
3  for i in range(20):
4      idx = i
5      print()
6      samples = FP['pad_data'].values[idx]
7      sample_rate = 16000
8      print('True Label      : ', FP['true_labels'][idx:idx+1].values[0])
9      print('Predicted label : ', FP['prediction'][idx:idx+1].values[0])
10
11     ipd.display(ipd.Audio(samples, rate = sample_rate))
12     print()
13
14
15     a = 211+i
16     fig = plt.figure(figsize=(14, 8))
17
18     plt.title('Raw wave of ' + FP['true_labels'][idx:idx+1].values[0])
19     plt.xlabel('Time(second)')
20     plt.ylabel('Amplitude')
21     plt.plot(np.linspace(0, sample_rate/len(samples), sample_rate), samples)
22     plt.show()

```

True Label : one

Predicted label : on

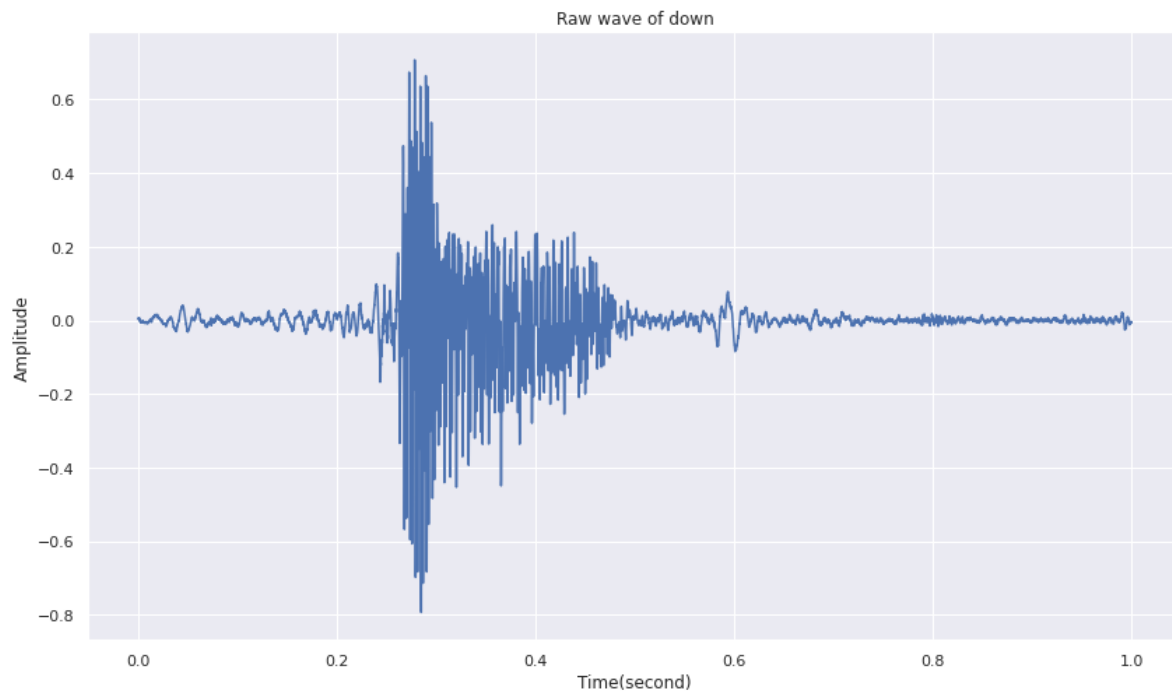
▶ 0:00 / 0:00  



True Label : down

Predicted label : dog

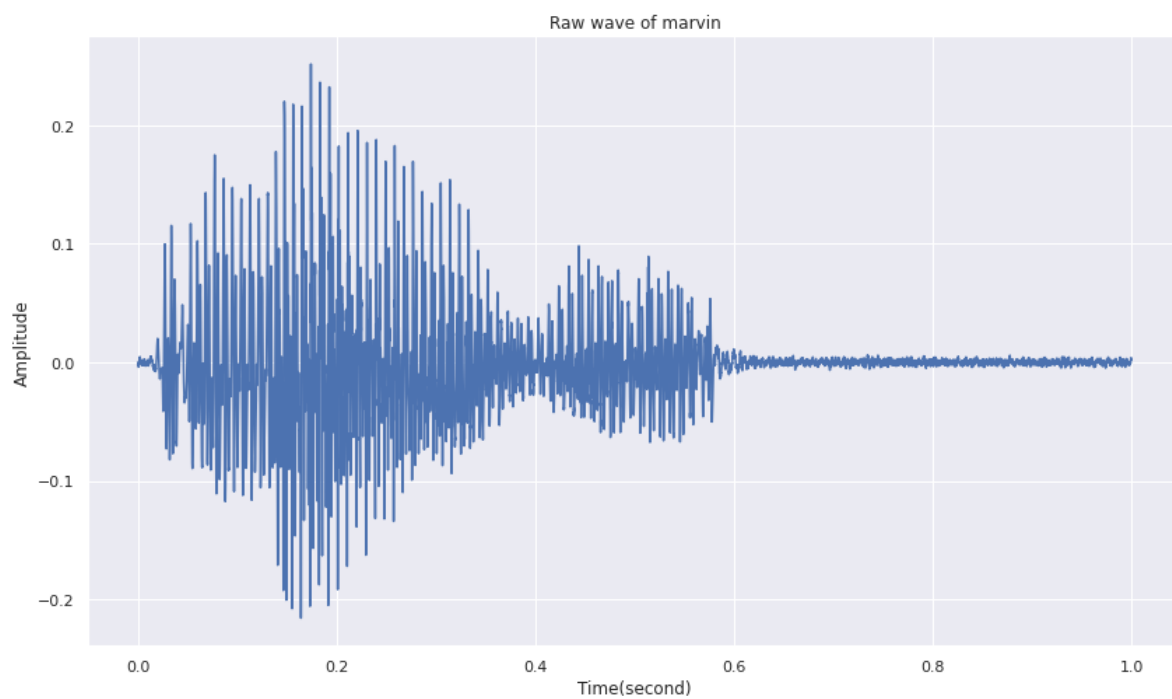
▶ 0:00 / 0:00



True Label : marvin

Predicted label : nine

▶ 0:00 / 0:00





True Label : cat

Predicted label : up

▶

0:00 / 0:00

🔊




True Label : marvin
Predicted label : nine

▶

0:00 / 0:00

🔊

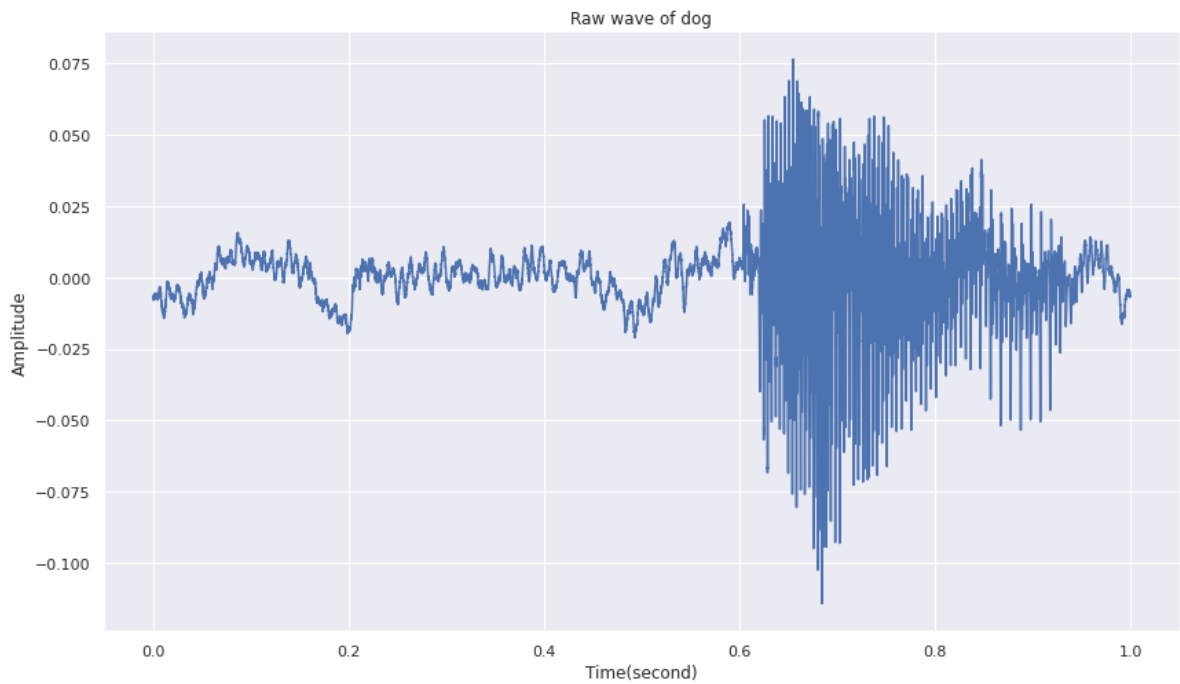



True Label : dog
Predicted label : down

▶

0:00 / 0:00

🔊



True Label : no
Predicted label : down

▶

0:00 / 0:00

🔊

Raw wave of no




True Label : tree
Predicted label : three

▶ 0:00 / 0:00



Raw wave of tree



True Label : wow
Predicted label : no

▶ 0:00 / 0:00



Raw wave of wow



True Label : stop
Predicted label : seven


▶

0:00 / 0:00

🔊



Raw wave of stop



True Label : on
Predicted label : one

▶

0:00 / 0:00

🔊

Raw wave of on




True Label : five
Predicted label : off

▶ 0:00 / 0:00 🔊



Raw wave of five



True Label : eight
Predicted label : right

▶ 0:00 / 0:00 🔊


Raw wave of eight



True Label : left
Predicted label : up


▶ 0:00 / 0:00 — 🔊

Raw wave of left




True Label : four
Predicted label : stop

▶ 0:00 / 0:00 — 🔊



True Label : no
Predicted label : right

▶ 0:00 / 0:00 — 🔊



True Label : go
Predicted label : no

▶ 0:00 / 0:00 — 🔊

Raw wave of go




True Label : go
Predicted label : right

▶

0:00 / 0:00

🔊

Raw wave of go



True Label : right
Predicted label : left

▶

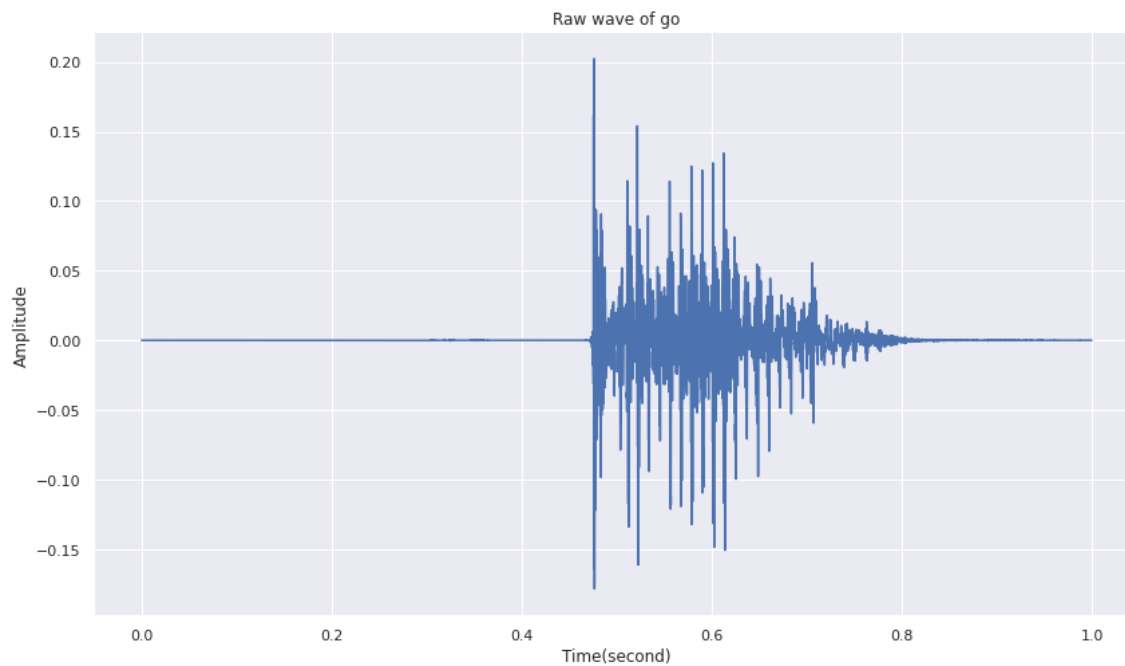
0:00 / 0:00

🔊



True Label : go
Predicted label : dog

▶ 0:00 / 0:00



Count plot of misclassified points

In []:

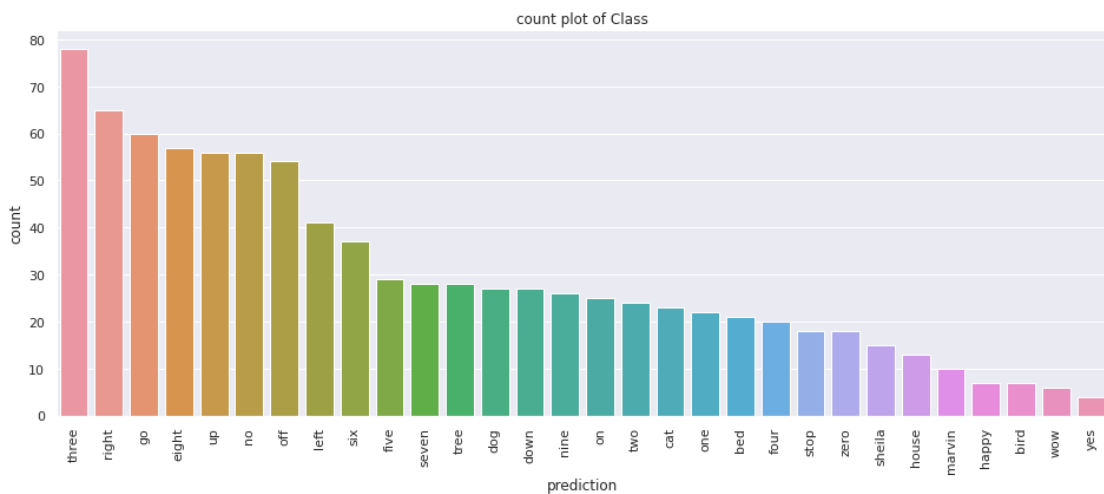
```

1 def plot(column_name , data, title):
2     #get unique_values of each categories
3     dic = Counter(data[column_name].values)
4     #sort in decending order by values
5     dic = sorted(dic.items(), key=lambda x: x[1], reverse=True)
6     column_list = []
7
8
9     #get name of sorted dic
10    for name in dic:
11        column_list.append(name[0])
12
13    plt.figure(figsize = (16,6))
14    ax = sns.countplot(x = column_name , data = data , order = column_list, dodge = F
15
16    h,l = ax.get_legend_handles_labels()
17    ax.legend(h ,column_list,bbox_to_anchor=(1.05, 1) ,loc = 'upper left')
18    plt.setp(ax.get_xticklabels() , rotation = 90 )
19    plt.title('count plot of {}'.format(title))
20    plt.show()
21    return dic

```

In []:

```
1 dic = plot('prediction' , FP, 'Class')
```



In []:

```

1 y_unique = list(set(list(FP['prediction'].values)))
2 len(y_unique)

```

Out[73]:

30

Classification_report

In []:

```
1 matrix = classification_report(df['true_labels'].values,df['prediction'].values)
```

In []:

```
1 print(matrix)
```

	precision	recall	f1-score	support
bed	0.94	0.93	0.94	343
bird	0.98	0.93	0.95	346
cat	0.94	0.96	0.95	347
dog	0.92	0.92	0.92	349
down	0.94	0.88	0.91	472
eight	0.89	0.96	0.92	470
five	0.94	0.91	0.92	471
four	0.96	0.96	0.96	475
go	0.88	0.91	0.89	474
happy	0.98	0.95	0.97	348
house	0.96	0.95	0.96	350
left	0.92	0.94	0.93	471
marvin	0.97	0.94	0.95	349
nine	0.94	0.93	0.94	473
no	0.89	0.94	0.91	475
off	0.89	0.93	0.91	471
on	0.94	0.86	0.90	473
one	0.95	0.94	0.95	474
right	0.87	0.95	0.91	473
seven	0.94	0.94	0.94	476
sheila	0.96	0.97	0.96	347
six	0.92	0.95	0.94	474
stop	0.96	0.92	0.94	476
three	0.84	0.90	0.87	471
tree	0.91	0.83	0.87	347
two	0.95	0.95	0.95	475
up	0.88	0.90	0.89	475
wow	0.98	0.96	0.97	349
yes	0.99	0.95	0.97	476
zero	0.96	0.94	0.95	475
accuracy			0.93	12945
macro avg	0.93	0.93	0.93	12945
weighted avg	0.93	0.93	0.93	12945

Confusion Matrix

In []:


```
1 #from sklearn.metrics import confusion_matrix
2 confusion = confusion_matrix(df['true_labels'].values, df['prediction'].values, labels
```

In []:

```

1 plt.figure(figsize = (35,35))
2 ax= plt.subplot()
3 sns.heatmap(confusion, annot=True, fmt='g', ax=ax); #annot=True to annotate cells, ft
4
5 # labels, title and ticks
6 ax.set_xlabel('Predicted labels');ax.set_ylabel('True labels');
7 ax.set_title('Confusion Matrix');
8 ax.xaxis.set_ticklabels(y_unique); ax.yaxis.set_ticklabels(y_unique);

```



Error Analysis

1. All misclassification points have very low amplitude.
2. Amplitude means : The amplitude of a wave is related to the amount of energy it carries. A low amplitude means wave of signal carries a small amount of energy.
3. Many True label are also labeled incorrectly. Ex: 3rd data point is labeled as marvie(Ture label) and predicted as nine and it is actual nine(concluded by listen). **Here many ground truth labels are wrong.** (cause of error : human error)
4. Many points are silent but labeled as it is not silent.
5. As we see in count plot highly misclassified label is three. and very low misclassified label is yes.
6. From confusion matrix we say that model is not able to classify tree as tree and three as three.
7. we can see in confusion matrix highest misclassification happen in three and tree(both are very similar in pronunciation).
8. From Classification Report we can conclude that,
 - **Model Perform Best**(where F1 score ≥ 95) in this categories : [bird, cat, four, happy, house, marvin, one, sheila, two, wow, yes, zero]
 - **Model Perform Medium**(where F1 score ≥ 90 and < 95) in this categories : [bed, dog, down, eight, five, left, nine, no, off, on, right, seven, six, stop]
 - **Model Perform Worst**(maximum misclassify, where F1 score < 91) in this categories : [go, three, tree, up]